

Avocado Price Analysis

February 15, 2021

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from scipy import stats
from scipy.stats import linregress
```

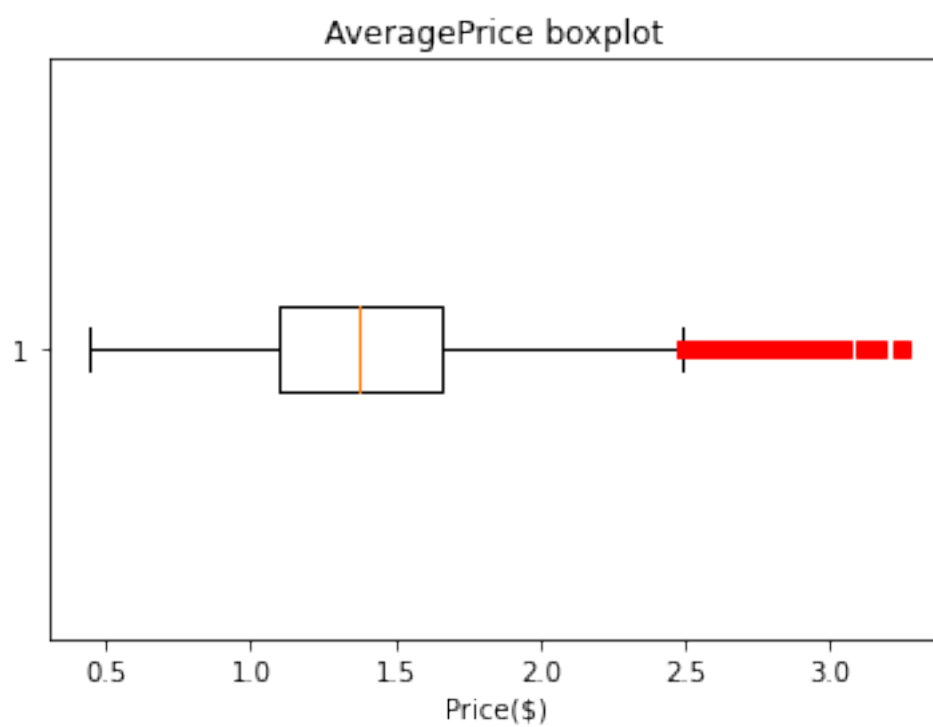
```
[2]: data = pd.read_csv('avocado.csv', delimiter = ',')
data.drop(columns='Unnamed: 0', inplace=True)
```

```
[3]: data.info()
```

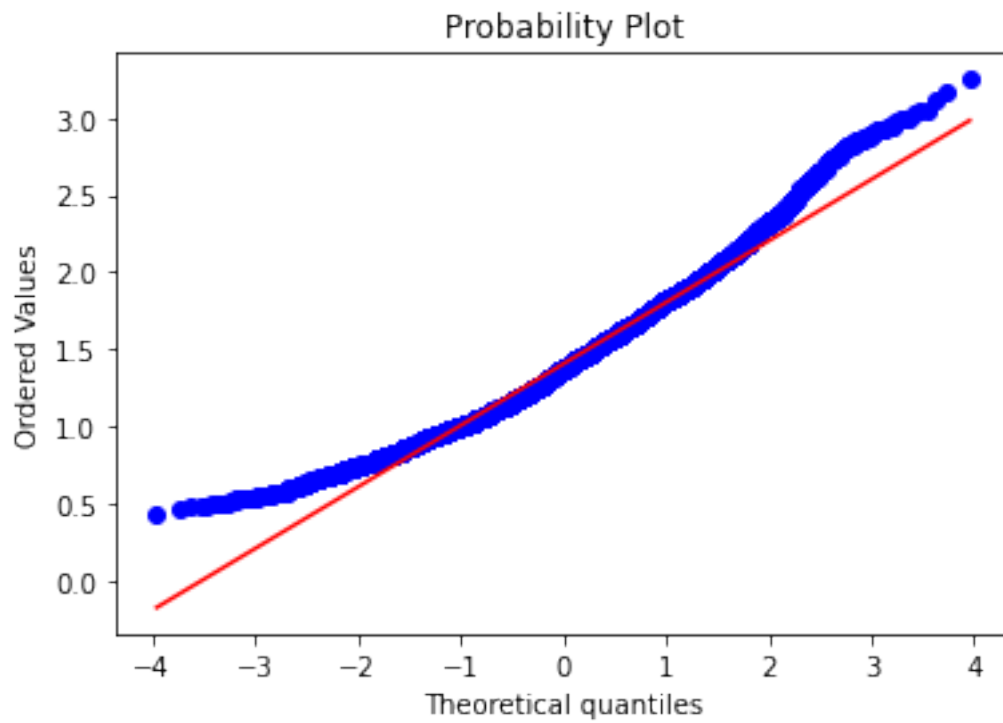
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18249 entries, 0 to 18248
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  18249 non-null  object
1   AveragePrice          18249 non-null  float64
2   Total Volume         18249 non-null  float64
3   4046                  18249 non-null  float64
4   4225                  18249 non-null  float64
5   4770                  18249 non-null  float64
6   Total Bags            18249 non-null  float64
7   Small Bags            18249 non-null  float64
8   Large Bags            18249 non-null  float64
9   XLarge Bags           18249 non-null  float64
10  type                  18249 non-null  object
11  year                  18249 non-null  int64
12  region                18249 non-null  object
dtypes: float64(9), int64(1), object(3)
memory usage: 1.8+ MB
```

```
[4]: plt.boxplot(data['AveragePrice'], 0, 'rs', 0)
plt.title('AveragePrice boxplot')
plt.xlabel('Price($)')
plt.plot()
```

[4]: []



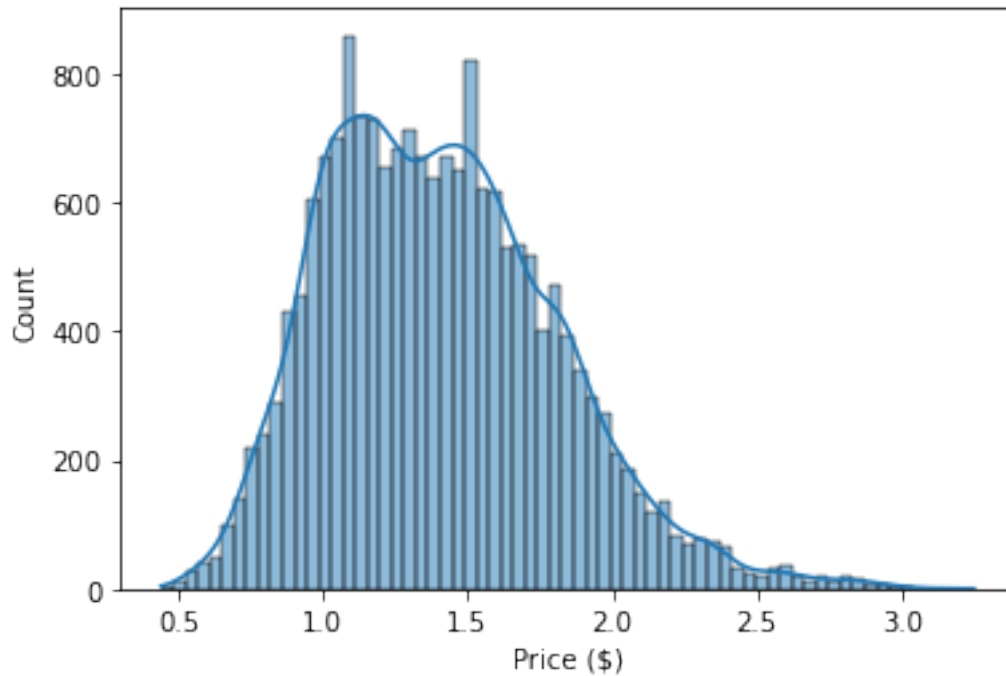
```
[5]: ax4 = plt.subplot()  
res = stats.probplot(data['AveragePrice'],dist = 'norm',plot=plt)
```



Not a great fit for a normal distriibution.

```
[6]: sns.histplot(data['AveragePrice'],kde=True)  
plt.xlabel('Price ($)')  
plt.plot()
```

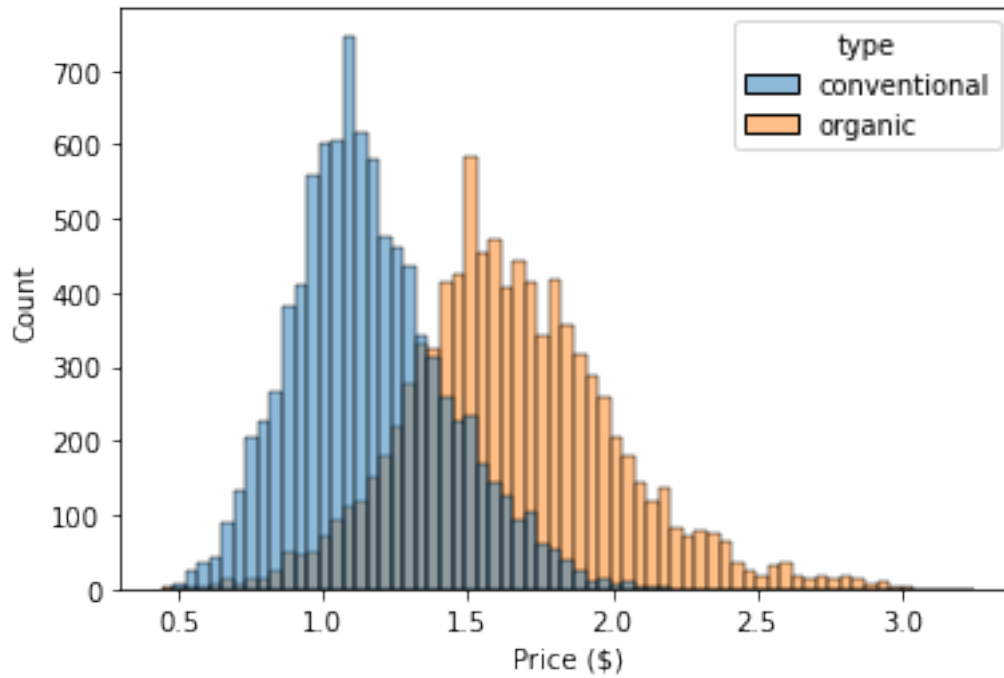
```
[6]: []
```



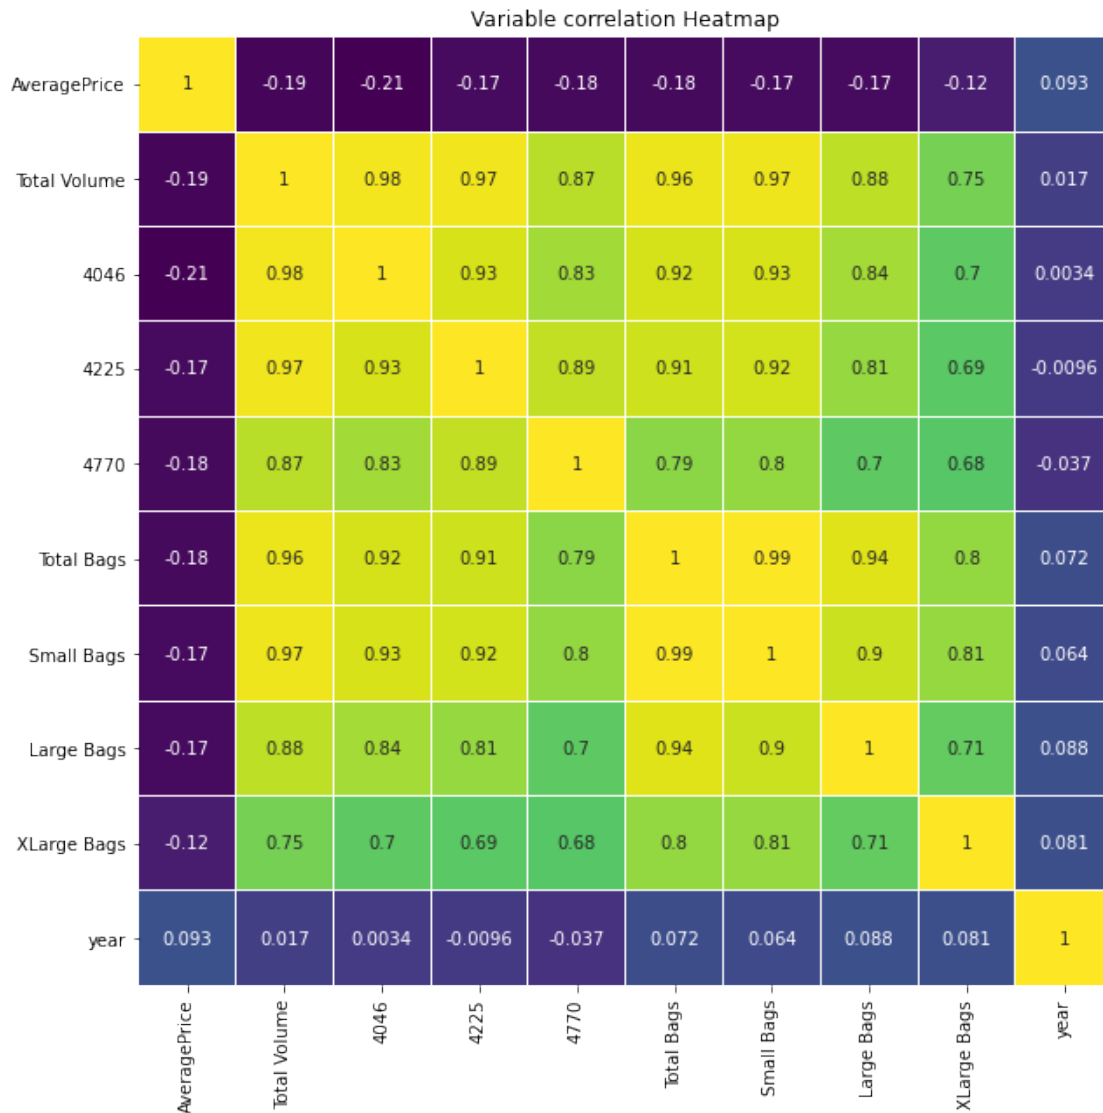
It appears to be a bimodal distribution, which is strange for a price of the same item you would expect it to be unimodal.

later in the analysis, I come to compare the variable 'type' which categorizes the avocados between organic and conventional. this explains the bimodal distribution, as the skews for organic and conv are different

```
[7]: sns.histplot(data,x='AveragePrice',hue='type')  
plt.xlabel('Price ($)')  
plt.show()
```



```
[8]: d_corr = data.corr()
fig, ax = plt.subplots(figsize=(10,10))
ax = sns.heatmap(d_corr, annot=True, linewidth=0.01, cbar=False, cmap='viridis')
ax.set_title('Variable correlation Heatmap')
plt.show()
```

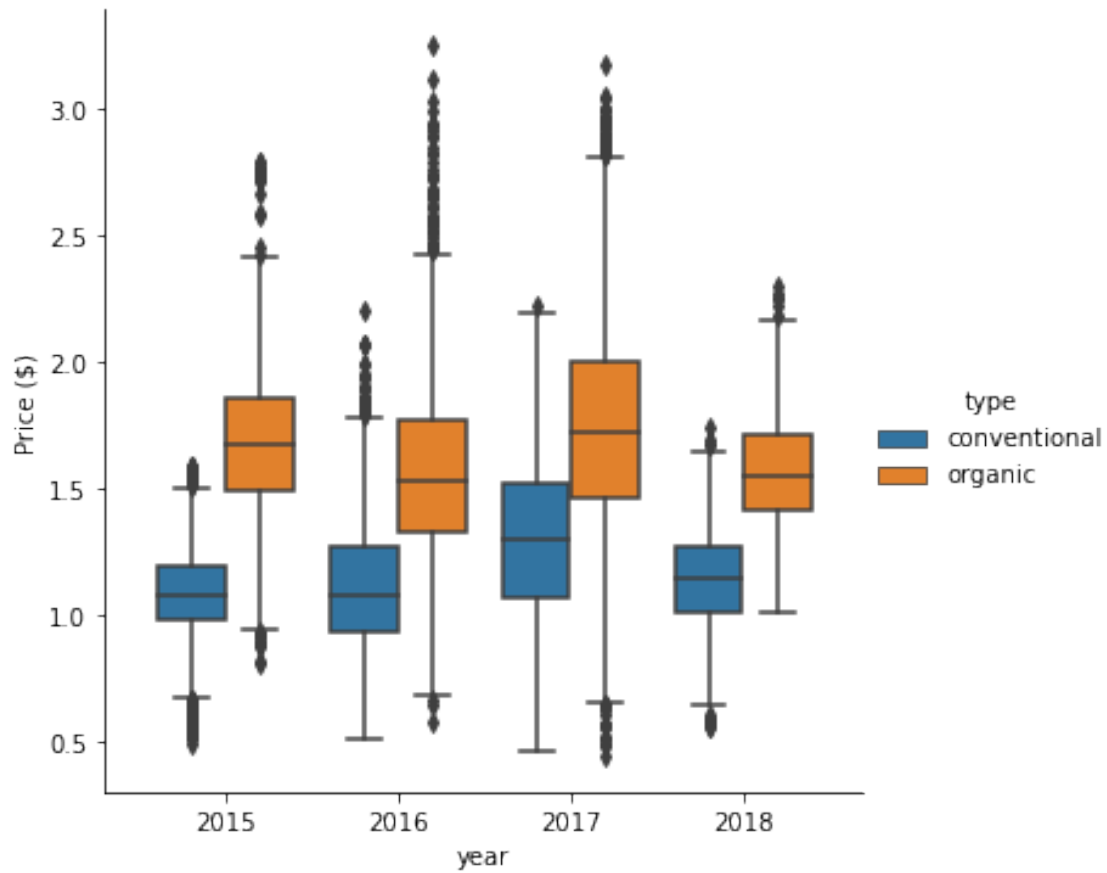


0.1 Categorical

0.1.1 Type organic / conventional

```
[9]: sns.catplot(x='year', y='AveragePrice', data=data, hue='type', kind='box')
plt.ylabel('Price ($)')
plt.plot()
```

```
[9]: []
```

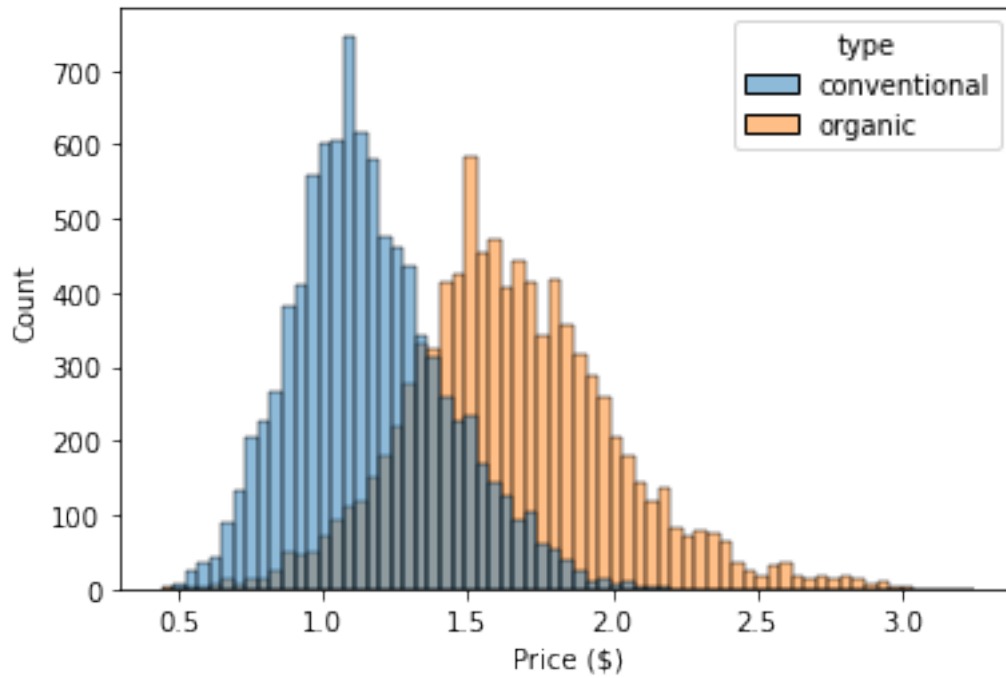


- Organic avocados tend to be more expensive than conventional avocados.

bernulli 1-organic, 0-conventional

```
[10]: sns.histplot(data, x='AveragePrice', hue='type')
plt.xlabel('Price ($)')
plt.plot()
```

[10]: []

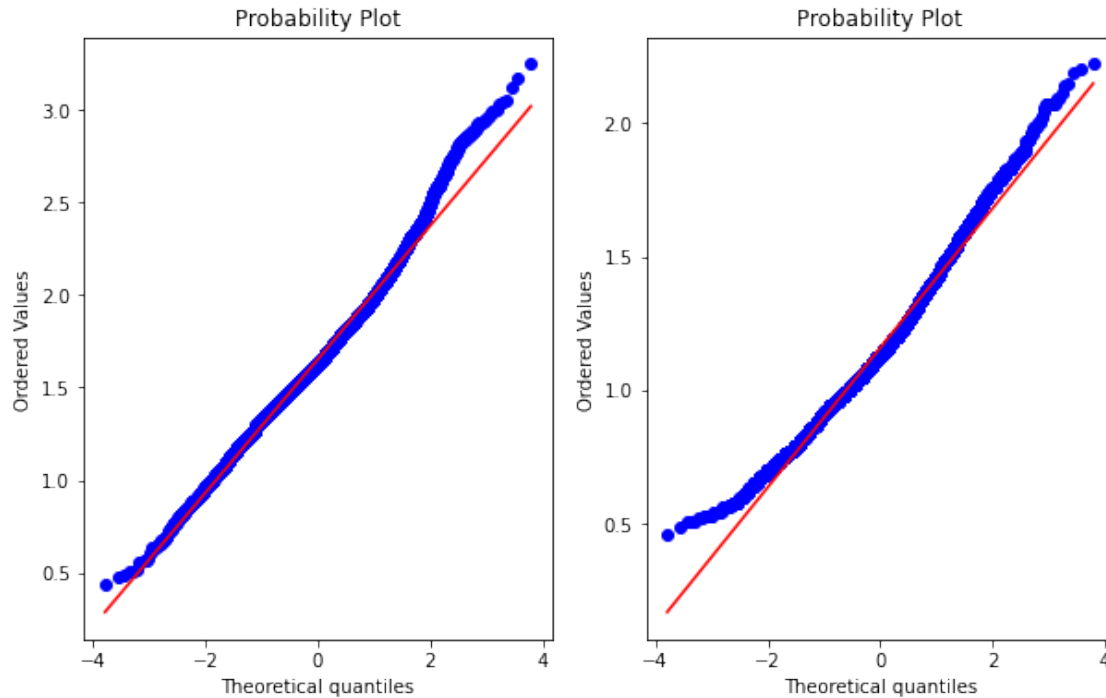


distribution fit

```
[11]: f, ax = plt.subplots(1,2,figsize=(10,6))

stats.probplot(data.loc[data['type']=='organic','AveragePrice'],dist_
↳='norm',plot=ax[0])
stats.probplot(data.loc[data['type']=='conventional','AveragePrice'],dist_
↳='norm',plot=ax[1])
plt.plot()
```

[11]: []



the spread of values for organic is much larger. $4Q > 3.0$ compared to $4Q < 3.0$

```
[12]: avo_type = data.groupby('type').agg(['mean', 'std', 'count'])
      avo_type = round(avo_type, 2)
      #mean diff
      mean_diff =
      ↪ avo_type['AveragePrice']['mean']['organic'] - avo_type['AveragePrice']['mean']['conventional']
      #standard deviation
      organic_s1 = avo_type['AveragePrice']['std']['organic']
      conventional_s2 = avo_type['AveragePrice']['std']['conventional']
      print('mean difference', mean_diff, 'p')
```

mean difference 0.49 p

X_1 = Organic Avocados, AveragePrice dist.

X_2 = Conventional Avocados, AveragePrice dist.

$\bar{X}_1 - \bar{X}_2 = 0.49$

$\alpha = 0.01$

$$\sigma_{\bar{X}_1 + \bar{X}_2}^2 \approx \frac{S_1}{n_1} + \frac{S_2}{n_2}$$

```
[13]: sampling_stderr = (organic_s1/
    ↪avo_type['AveragePrice']['count']['organic'])+(conventional_s2/
    ↪avo_type['AveragePrice']['count']['conventional'])
sampling_stderr = np.sqrt(sampling_stderr)
print('sampling std error approx',sampling_stderr)

crit_limit =2.33
limit = crit_limit*sampling_stderr
print(round(mean_diff - limit,2),'to',round(mean_diff +limit,2))
```

sampling std error approx 0.008243223411136655
0.47 to 0.51

Confident (That the true Mean difference in price between Organic avocados and conventional avocados is between 0.47p and 0.51p) $\approx 99\%$

0.2 Region

Groupby region, with lambda agg to reduce outlier prices.

```
[14]: region_data = data.groupby(['region','type'])['AveragePrice'].agg(['mean',
    ↪                                lambda x : np.
    ↪                                quantile(x,.15),
    ↪                                lambda x : np.
    ↪                                quantile(x,.85)])
region_data = region_data.unstack()
region_data.rename(columns={'<lambda_0>':'q15','<lambda_1>':'q85'},inplace=True)
err_pdata = region_data.copy()
region_data.head()
```

```
[14]:
```

	mean		q15		q85 \
type	conventional	organic	conventional	organic	conventional
region					
Albany	1.348757	1.773314	1.110	1.54	1.588
Atlanta	1.068817	1.607101	0.902	1.25	1.230
BaltimoreWashington	1.344201	1.724260	1.120	1.51	1.600
Boise	1.076036	1.620237	0.836	1.16	1.268
Boston	1.304379	1.757396	1.070	1.49	1.580

type	organic
region	
Albany	1.970
Atlanta	1.906
BaltimoreWashington	2.028
Boise	2.088
Boston	2.010

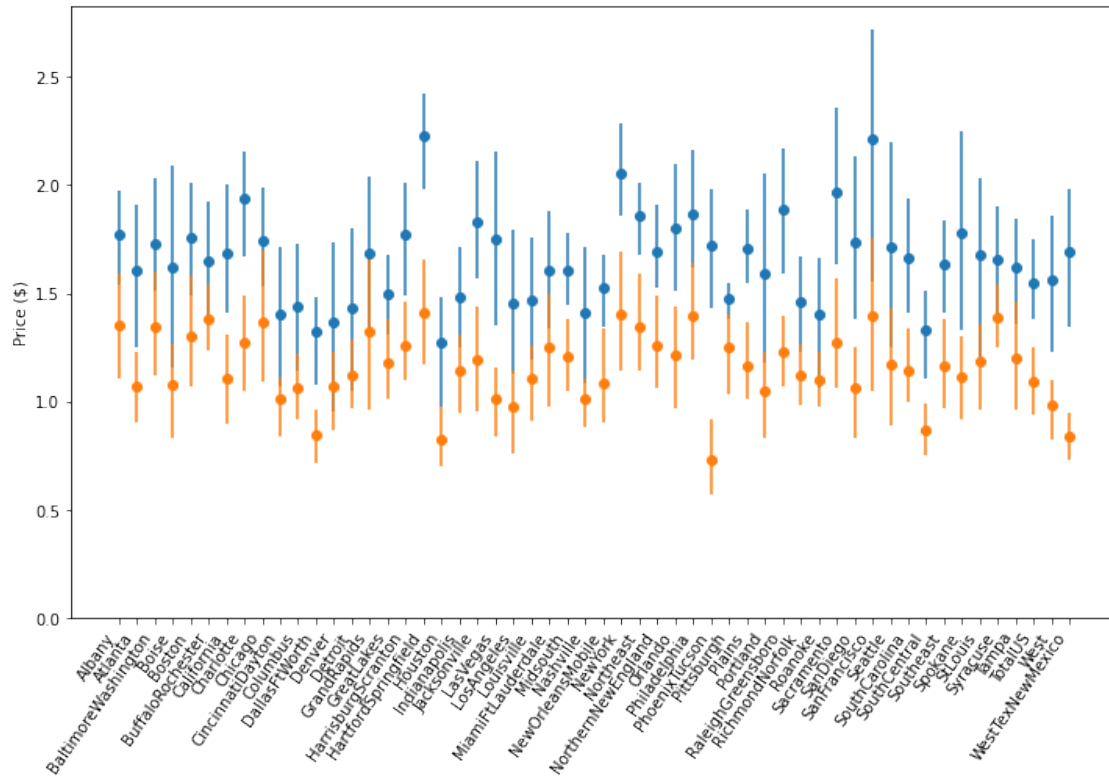
```
[15]: #reformatting the quartile columns for the error bar plot
err_pdata['q15'] = err_pdata['mean'] - err_pdata['q15']
err_pdata['q85'] = err_pdata['q85'] - err_pdata['mean']
err_pdata.head()
```

```
[15]:
```

	mean		q15		
type	conventional	organic	conventional	organic	\
region					
Albany	1.348757	1.773314	0.238757	0.233314	
Atlanta	1.068817	1.607101	0.166817	0.357101	
BaltimoreWashington	1.344201	1.724260	0.224201	0.214260	
Boise	1.076036	1.620237	0.240036	0.460237	
Boston	1.304379	1.757396	0.234379	0.267396	

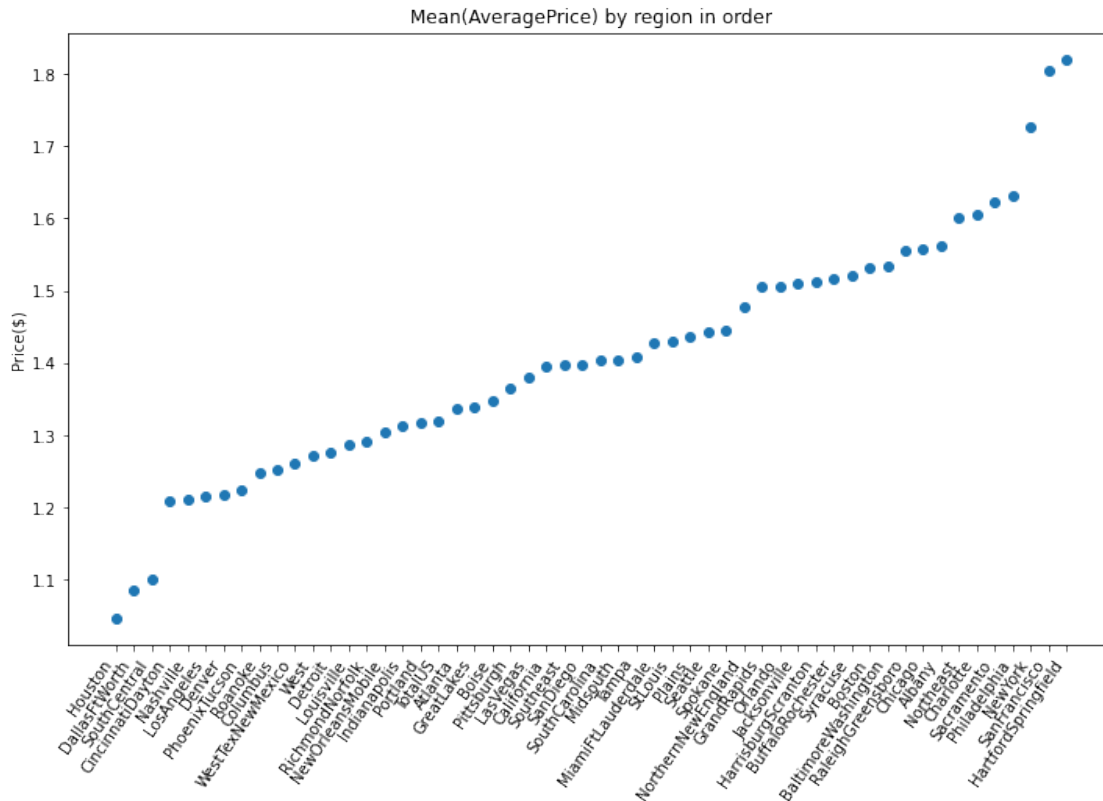
	q85		
type	conventional	organic	
region			
Albany	0.239243	0.196686	
Atlanta	0.161183	0.298899	
BaltimoreWashington	0.255799	0.303740	
Boise	0.191964	0.467763	
Boston	0.275621	0.252604	

```
[16]: fig,ax = plt.subplots(figsize=(12,8))
fig.autofmt_xdate(rotation=55 )
ax.errorbar(x = err_pdata.
↳index,y=err_pdata['mean']['organic'],yerr=[err_pdata['q15']['organic'],err_pdata['q85']['or
↳='o')
ax.errorbar(x = err_pdata.
↳index,y=err_pdata['mean']['conventional'],yerr=[err_pdata['q15']['conventional'],err_pdata[
↳='o')
ax.set_yticks([0,0.5,1,1.5,2,2.5])
ax.set_ylabel('Price ($)')
plt.show()
```



```
[17]: price_by_region = data.groupby('region')['AveragePrice'].agg(['mean'])
price_by_region = price_by_region.sort_values('mean')
price_by_region = price_by_region.reset_index()
fig,ax = plt.subplots(figsize = (12,8))
fig.autofmt_xdate(rotation=55 )
ax.scatter(x=price_by_region['region'],y=price_by_region['mean'])
ax.set_title('Mean(AveragePrice) by region in order')
ax.set_ylabel('Price($)')
plt.plot()
```

[17]: []



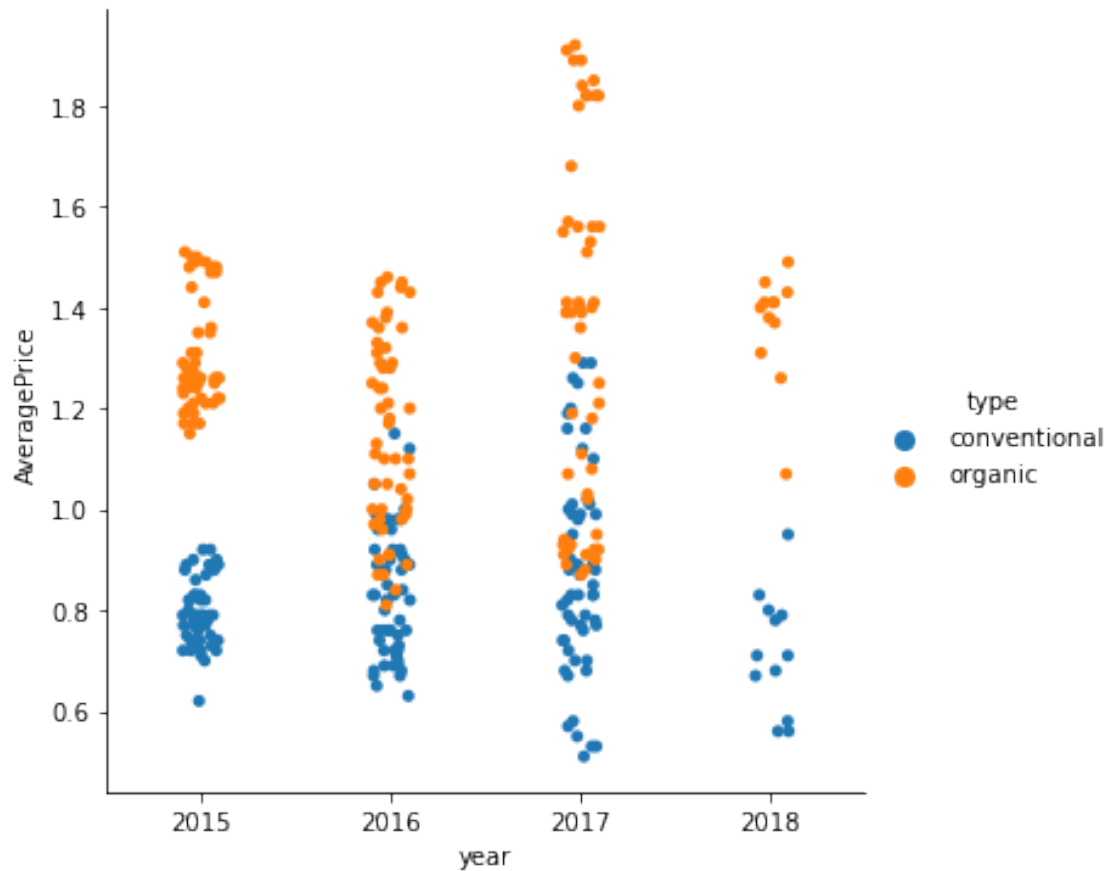
- Houston on average has the cheapest avocados
- Houston also has on average the cheapest organic avocados
- pittsburg on average has the cheapest Conventional avocados

Even though i think i will be a terrible fit, i want to fit a linear regression line to Price over year.

```
[18]: houston = data[data['region']=='Houston']
      h_organic = houston[houston['type']=='organic']
      h_conv = houston[houston['type']=='conventional']
```

```
[19]: sns.catplot(data = houston,x='year',y='AveragePrice',hue='type')
```

```
[19]: <seaborn.axisgrid.FacetGrid at 0x16ef9db6f08>
```

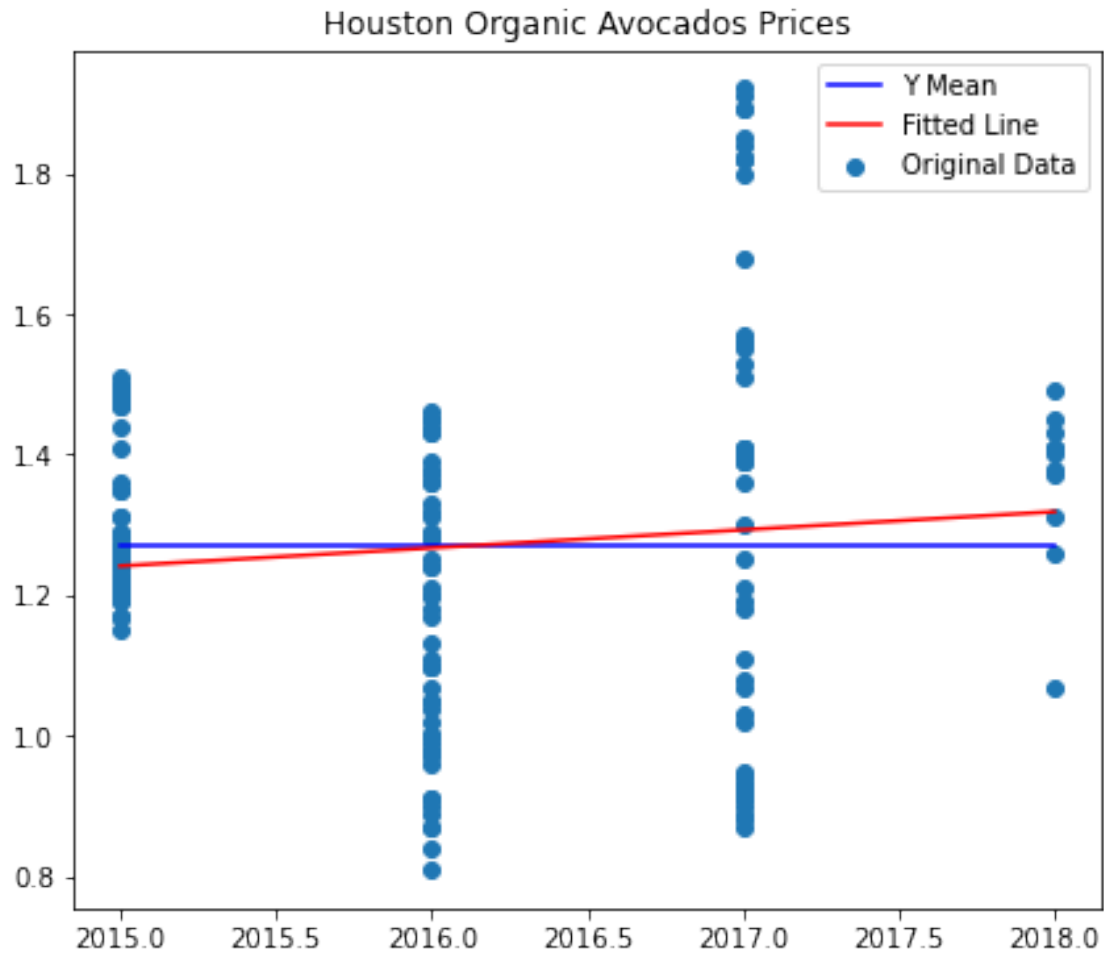


0.2.1 Organic

```
[20]: slope_org, intercept_org, r_org, p_org, se_org =  $\square$ 
      \(\rightarrow\)linregress(h_organic['year'],h_organic['AveragePrice'])
```

```
[21]: fig = plt.figure(figsize=(7,6))
plt.scatter(h_organic['year'],h_organic['AveragePrice'],label='Original Data')
plt.plot(h_organic['year'],h_organic['year'].apply(lambda val :
\(\rightarrow\)h_organic['AveragePrice'].mean()),label = 'Y Mean',c='b')
plt.plot(h_organic['year'],h_organic['year'].apply(lambda val : val*slope_org +  $\square$ 
\(\rightarrow\)intercept_org),label='Fitted Line',c='r')
plt.title('Houston Organic Avocados Prices')
plt.legend()
plt.plot()
```

```
[21]: []
```



coefficient of determination

```
[22]: h_organic = h_organic.copy()
      h_organic.loc[:, 'y'] = h_organic.loc[:, 'AveragePrice']
      h_organic.loc[:, 'mx+b'] = h_organic.loc[:, 'year'].apply(lambda val :
      ↪ val*slope_org + intercept_org)
      h_organic.loc[:, 'SE_line'] = (h_organic.loc[:, 'y'] - h_organic.loc[:,
      ↪ 'mx+b'])**2

[23]: one_minusr2 = round((1- h_organic.loc[:, 'SE_line'].sum()/((h_organic.loc[:,
      ↪ 'y']-h_organic.loc[:, 'y'].mean())**2).sum()),2)
      print(one_minusr2*100, '% of the total variation is described by the regression_
      ↪ line y=', round(slope_org,2), 'x', round(intercept_org,2))
```

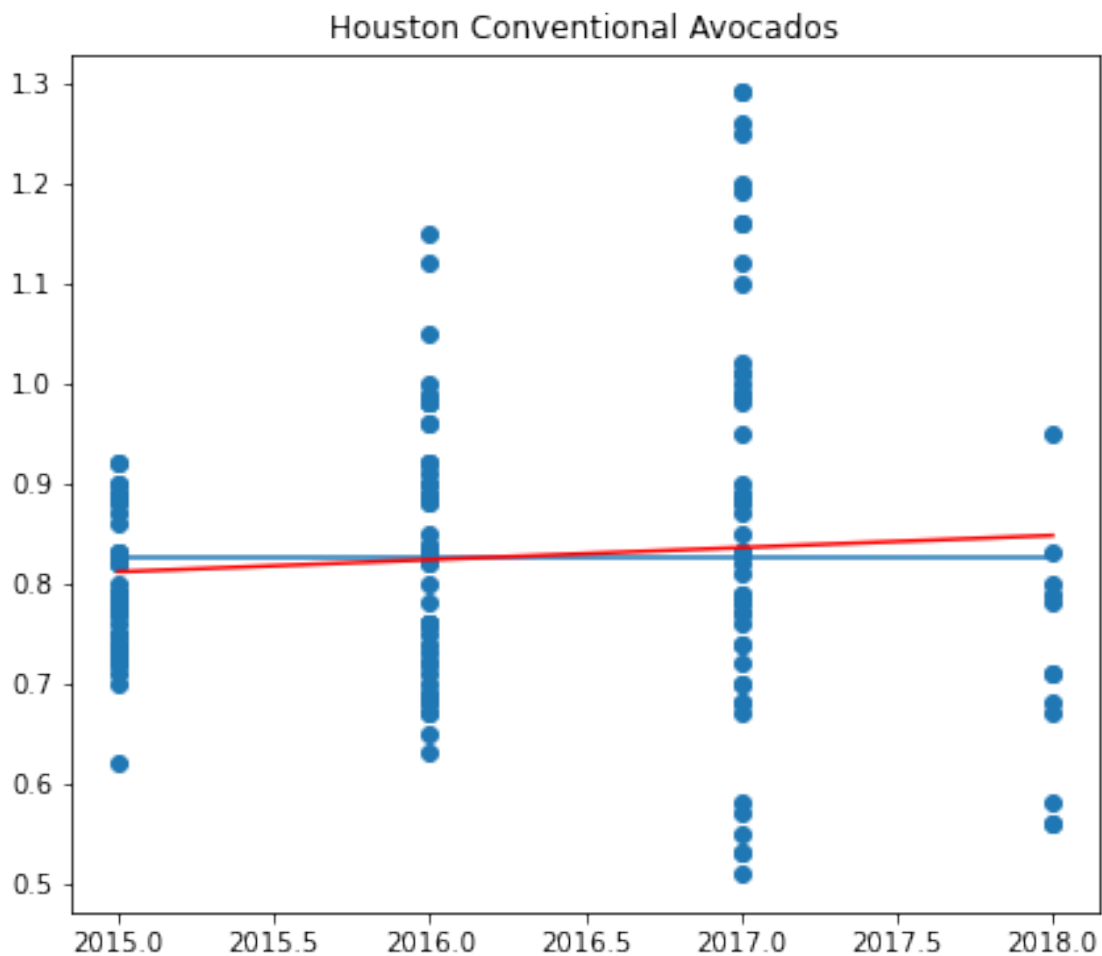
1.0 % of the total variation is described by the regression line $y = 0.03x - 50.59$

0.2.2 Conventional

```
[24]: slope_con, intercept_con, r, p, se =  $\hookrightarrow$ linregress(h_conv['year'],h_conv['AveragePrice'])
```

```
[25]: fig = plt.figure(figsize=(7,6))  
plt.scatter(h_conv['year'],h_conv['AveragePrice'])  
plt.plot(h_conv['year'],h_conv['year'].apply( $\hookrightarrow$ lambda val : h_conv['AveragePrice'].  
     $\hookrightarrow$ mean()))  
plt.plot(h_conv['year'],h_conv['year'].apply( $\hookrightarrow$ lambda val : val*slope_con +  $\hookrightarrow$   
     $\hookrightarrow$ intercept_con),c='r')  
plt.title('Houston Conventional Avocados')  
plt.plot()
```

```
[25]: []
```



coefficient of determination


```
[26]: h_conv = h_conv.copy()
h_conv.loc[:, 'y'] = h_conv.loc[:, 'AveragePrice']
h_conv.loc[:, 'mx+b'] = h_conv.loc[:, 'year'].apply(lambda val : val*slope_con +
↪ intercept_con)
h_conv.loc[:, 'SE_line'] = (h_conv.loc[:, 'y'] - h_conv.loc[:, 'mx+b'])**2
```

```
[27]: one_minusr2_con = round((1- h_conv.loc[:, 'SE_line'].sum()/((h_conv.loc[:,
↪ 'y']-h_conv.loc[:, 'y'].mean())**2).sum()),2)
print(one_minusr2_con*100,'% of the total variation is described by the
↪ regression line y=',round(slope_con,2),'x',round(intercept_con,2))
```

1.0 % of the total variation is described by the regression line $y = 0.01 x - 23.65$

- Not surprisingly simple linear regression, least squared line has not produced a reliable predictor, non the the less i wanted to put it into practive and visualize it.

```
[28]: #noramlizing
from sklearn import preprocessing
from sklearn.decomposition import PCA
from sklearn.linear_model import LinearRegression
```

```
[29]: def standardize(column,int_= False):
    if int_ ==True:
        values = normalized[column].sort_values().unique()
    else:
        values = normalized[column].unique()

    for i,v in enumerate(values):
        normalized.loc[normalized[column]==v, column]=i
```

```
[30]: normalized = data.copy()
```

```
[31]: #type- binary
normalized.loc[normalized['type'] == 'conventional','type']=0
normalized.loc[normalized['type'] == 'organic','type']=1

#region -nominal
standardize('region')

#year - ordinal
standardize('year',True)

#Date - ordinal
standardize('Date',True)
```

```
[32]: normalized.head()
```

```
[32]:
```

	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags \
0	51	1.33	64236.62	1036.74	54454.85	48.16	8696.87
1	50	1.35	54876.98	674.28	44638.81	58.33	9505.56
2	49	0.93	118220.22	794.70	109149.67	130.50	8145.35
3	48	1.08	78992.15	1132.00	71976.41	72.58	5811.16
4	47	1.28	51039.60	941.48	43838.39	75.78	6183.95

	Small Bags	Large Bags	XLarge Bags	type	year	region
0	8603.62	93.25	0.0	0	0	0
1	9408.07	97.49	0.0	0	0	0
2	8042.21	103.14	0.0	0	0	0
3	5677.40	133.76	0.0	0	0	0
4	5986.26	197.69	0.0	0	0	0

```
[33]: #dir(preprocessing)
```

```
[34]: scaler = preprocessing.MinMaxScaler(feature_range = (-1,1))
names = normalized.columns

d = scaler.fit_transform(normalized)

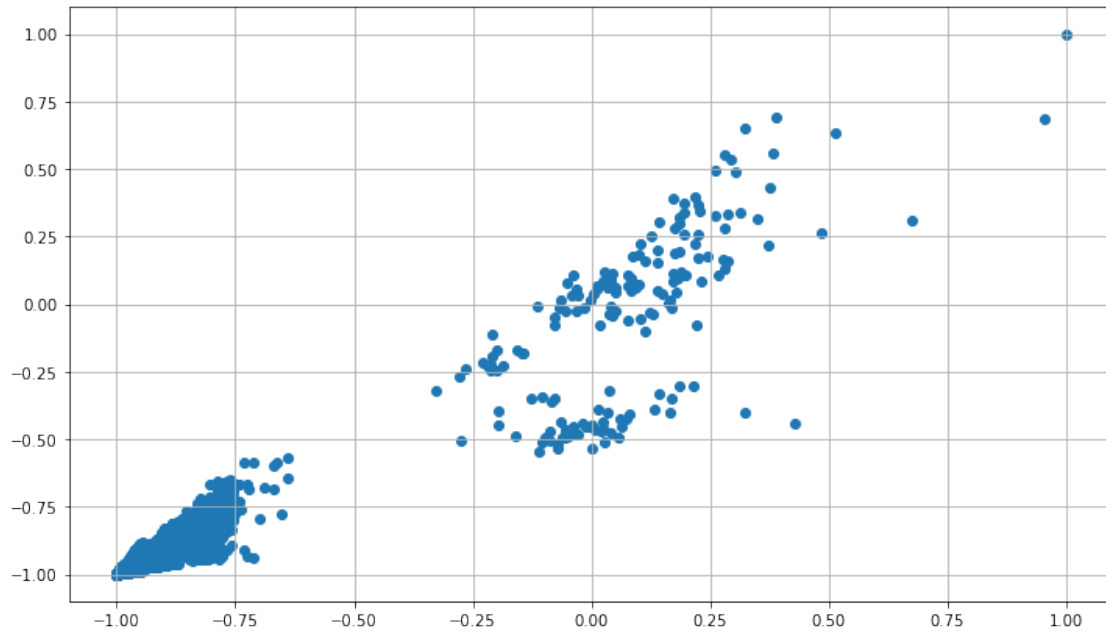
stand_data=pd.DataFrame(data=d, columns=names)
stand_data.head()
```

```
[34]:
```

	Date	AveragePrice	Total Volume	4046	4225	4770	\
0	-0.392857	-0.366548	-0.997947	-0.999909	-0.994680	-0.999962	
1	-0.404762	-0.352313	-0.998247	-0.999941	-0.995639	-0.999954	
2	-0.416667	-0.651246	-0.996220	-0.999930	-0.989336	-0.999898	
3	-0.428571	-0.544484	-0.997475	-0.999900	-0.992968	-0.999943	
4	-0.440476	-0.402135	-0.998370	-0.999917	-0.995717	-0.999940	

	Total Bags	Small Bags	Large Bags	XLarge Bags	type	year	region
0	-0.999102	-0.998714	-0.999967	-1.0	-1.0	-1.0	-1.0
1	-0.999019	-0.998594	-0.999966	-1.0	-1.0	-1.0	-1.0
2	-0.999159	-0.998798	-0.999964	-1.0	-1.0	-1.0	-1.0
3	-0.999400	-0.999152	-0.999953	-1.0	-1.0	-1.0	-1.0
4	-0.999362	-0.999105	-0.999931	-1.0	-1.0	-1.0	-1.0

```
[35]: plt.figure(figsize=(12,7))
plt.scatter(stand_data['Total Volume'],stand_data['Total Bags'])
plt.grid()
```



```
[36]: correlations = stand_data.corr()
correlations[(correlations>0.5) | (correlations <-0.5)]
```

```
[36]:
```

	Date	AveragePrice	Total Volume	4046	4225	\
Date	1.000000	NaN	NaN	NaN	NaN	
AveragePrice	NaN	1.000000	NaN	NaN	NaN	
Total Volume	NaN	NaN	1.000000	0.977863	0.974181	
4046	NaN	NaN	0.977863	1.000000	0.926110	
4225	NaN	NaN	0.974181	0.926110	1.000000	
4770	NaN	NaN	0.872202	0.833389	0.887855	
Total Bags	NaN	NaN	0.963047	0.920057	0.905787	
Small Bags	NaN	NaN	0.967238	0.925280	0.916031	
Large Bags	NaN	NaN	0.880640	0.838645	0.810015	
XLarge Bags	NaN	NaN	0.747157	0.699377	0.688809	
type	NaN	0.615845	NaN	NaN	NaN	
year	0.950274	NaN	NaN	NaN	NaN	
region	NaN	NaN	NaN	NaN	NaN	

	4770	Total Bags	Small Bags	Large Bags	XLarge Bags	\
Date	NaN	NaN	NaN	NaN	NaN	
AveragePrice	NaN	NaN	NaN	NaN	NaN	
Total Volume	0.872202	0.963047	0.967238	0.880640	0.747157	
4046	0.833389	0.920057	0.925280	0.838645	0.699377	
4225	0.887855	0.905787	0.916031	0.810015	0.688809	
4770	1.000000	0.792314	0.802733	0.698471	0.679861	
Total Bags	0.792314	1.000000	0.994335	0.943009	0.804233	

Small Bags	0.802733	0.994335	1.000000	0.902589	0.806845
Large Bags	0.698471	0.943009	0.902589	1.000000	0.710858
XLarge Bags	0.679861	0.804233	0.806845	0.710858	1.000000
type	NaN	NaN	NaN	NaN	NaN
year	NaN	NaN	NaN	NaN	NaN
region	NaN	NaN	NaN	NaN	NaN

	type	year	region
Date	NaN	0.950274	NaN
AveragePrice	0.615845	NaN	NaN
Total Volume	NaN	NaN	NaN
4046	NaN	NaN	NaN
4225	NaN	NaN	NaN
4770	NaN	NaN	NaN
Total Bags	NaN	NaN	NaN
Small Bags	NaN	NaN	NaN
Large Bags	NaN	NaN	NaN
XLarge Bags	NaN	NaN	NaN
type	1.000000	NaN	NaN
year	NaN	1.000000	NaN
region	NaN	NaN	1.0

0.3 PCA

which attributes account for the least variation on our data, those will be the least necessary

```
[62]: pca = PCA(n_components=stand_data.shape[1])
```

```
pca_fitted = pca.fit(stand_data)
components = pca.transform(stand_data)
```

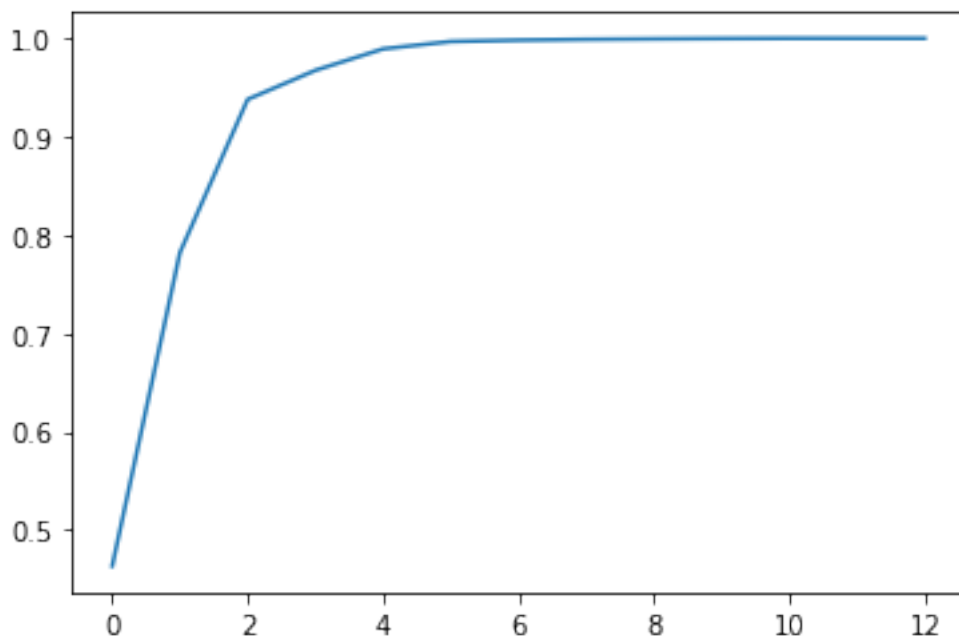
```
[63]: #this is our ratio of variance explained by attributes.
```

```
var = pca_fitted.explained_variance_ratio_
```

```
#now we can visually see the redundant attributes
```

```
plt.plot(np.cumsum(var))
plt.plot()
```

```
[63]: []
```



```
[64]: pca_c = ['pca_'+str(i) for i in range(0,components.shape[1])]
      pca_data = pd.DataFrame(components,columns=pca_c)
      pca_data
```

```
[64]:
```

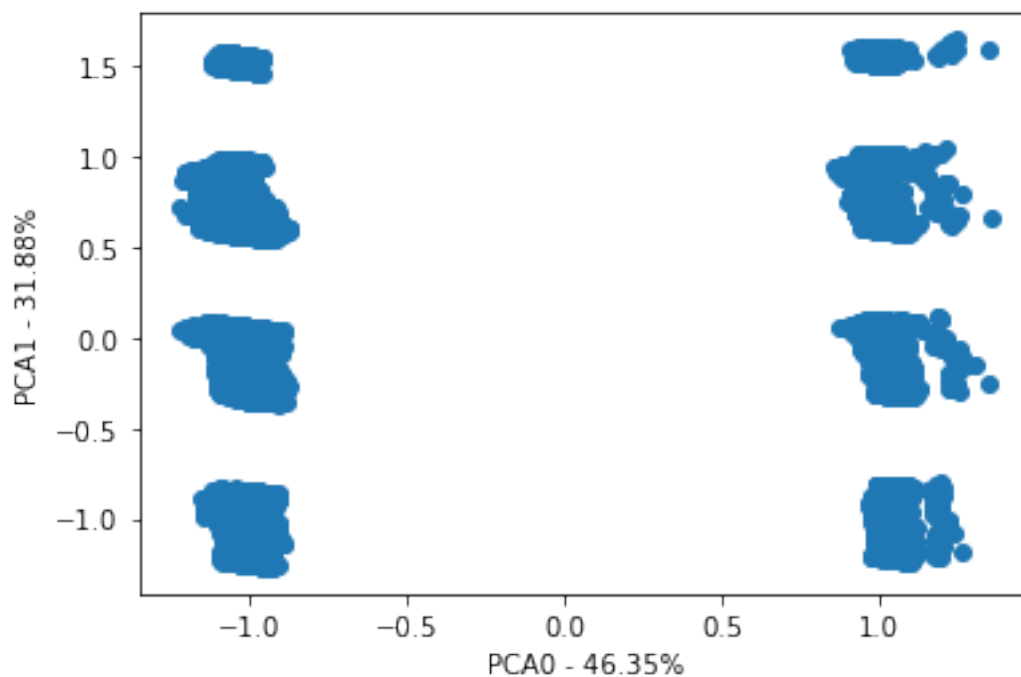
	pca_0	pca_1	pca_2	pca_3	pca_4	pca_5	pca_6 \
0	0.995166	-0.809407	-1.010630	-0.061394	0.191101	0.196340	-0.016465
1	0.992663	-0.816917	-1.010737	-0.063891	0.203148	0.184953	-0.017301
2	1.047656	-0.836696	-1.008749	-0.009316	-0.080734	0.234164	-0.012389
3	1.028121	-0.840589	-1.009538	-0.029344	0.018608	0.204874	-0.014889
4	1.002134	-0.843083	-1.010455	-0.054671	0.151833	0.168688	-0.017434
...
18244	-1.033556	1.511647	0.997426	-0.084335	-0.132351	-0.146628	0.027066
18245	-1.043858	1.505797	0.997079	-0.094067	-0.079813	-0.166299	0.026818
18246	-1.064532	1.502185	0.996460	-0.112975	0.026778	-0.196996	0.026047
18247	-1.072192	1.495792	0.996255	-0.119661	0.065945	-0.213902	0.025417
18248	-1.031736	1.479087	0.997513	-0.081675	-0.144471	-0.179802	0.026485
...
	pca_7	pca_8	pca_9	pca_10	pca_11	pca_12	
0	-0.010422	0.000401	-0.003641	0.001072	7.966001e-08	1.284986e-10	
1	-0.010468	0.000321	-0.004473	0.000530	8.228737e-08	1.237921e-10	
2	-0.007996	0.000881	0.005166	0.004296	5.861508e-08	1.297201e-10	
3	-0.008751	0.000582	0.000689	0.002713	7.046306e-08	1.237488e-10	
4	-0.010013	0.000348	-0.003803	0.001082	8.515053e-08	1.129118e-10	
...	
18244	0.010713	-0.001993	0.003924	-0.000413	-2.549018e-07	6.141789e-11	
18245	0.010031	-0.001929	0.002736	-0.000041	-2.453427e-07	5.232955e-11	

```
18246  0.008989 -0.001803  0.000255 -0.000658 -2.323529e-07  3.902631e-11
18247  0.008436 -0.001870 -0.000510 -0.000927 -2.265541e-07  3.037124e-11
18248  0.010710 -0.002090  0.003911 -0.000270 -2.466722e-07  4.163929e-11
```

[18249 rows x 13 columns]

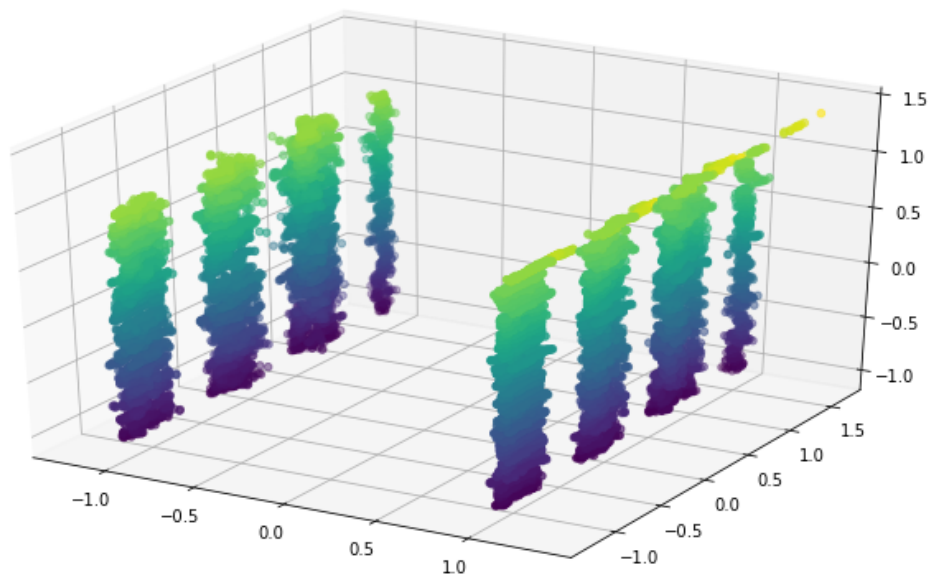
```
[65]: plt.scatter(pca_data.pca_0,pca_data.pca_1)
plt.xlabel('PCA0 - {0}%'.format(round(var[0]*100,2)))
plt.ylabel('PCA1 - {0}%'.format(round(var[1]*100,2)))
plt.plot()
```

[65]: []



```
[66]: plt.figure(figsize=(12,7))

ax = plt.axes(projection='3d')
xdata = pca_data.pca_0
ydata = pca_data.pca_1
zdata = pca_data.pca_2
ax.scatter3D(xdata,ydata,zdata,c=zdata)
plt.show()
```



[]: